

# kayak\_explore\_scraper

September 17, 2019

The Kayak explore feature uses google maps to display the cheapest flights to international airports across the world at any point within a specified time interval. You can leave the interval blank to find the cheapest flights possible, but in this case we are looking for a good deal within the next summer break (specifically, the beginning of June to the end of August). Our program first uses the python request library to scrape all of the flight data sent to google maps from kayak in JSON format. We then parse the JSON string to get the specific details we are interested in.

```
In [1]: import requests, smtplib, os, datetime
import pandas as pd
from bs4 import *
import urllib.request as ur
from email.mime.multipart import MIMEMultipart
from email.mime.text import MIMEText
from matplotlib import pyplot as plt

# Specify the beginning and end of the time frame of possible dates as YYYYMMDD
timeframe_begin = 20200601
timeframe_end = 20200830

def scrape_kayak(start='', end='', airport = 'BER'):
    """
    This function scrapes flight information from the kayak explore page.

    Parameters:
    start, end, airport - integer representing earliest possible departure date
    in YYYYMMDD format, integer representing latest return date, string with
    three letter code for starting airport. When both are start and end are
    left blank, results are returned from present date to one year in the
    future.

    Returns:
    df - a data frame containing all destination cities and corresponding
    flight information returned by the scrape
    """

    # Format the beginning and end dates to insert them into the URL
    start = '&depart=' + str(start)
```

```

end = '&return=' + str(end)

url = "https://www.kayak.com/s/horizon/exploreapi/elasticbox?airport=" + airport +
"&stopsFilterActive=false&duration=&budget=&topRightLat=68.58212830775821&topRightLat=" +
response = requests.post(url).json()

df = pd.DataFrame(columns=['City', 'Country', 'Price', 'Airline', 'Airport', 'Date'])

for i in range(len(response['destinations'])):
    destination = response['destinations'][i]
    row = list([destination['city']['name'], destination['country']['name'],
                destination['flightInfo']['price'], destination['airline'],
                destination['airport']['shortName'], pd.to_datetime(destination['departure']),
                str('http://kayak.com'+destination['clickoutUrl'])])
    df.loc[i] = row

city_mins = df.groupby(['City']).idxmin().astype(int)
df = df.loc[city_mins['Price'].to_list()]
# There is a glitch where some flights are returned with unrealistically
# high prices, so we'll remove those entries.
df = df.where(df['Price']!=999999).dropna()

return df

all_flights = scrape_kayak(timeframe_begin, timeframe_end)
all_flights.head()

```

```

Out[1]:

```

	City	Country	Price	Airline	Airport	\
295	Abbotsford	Canada	1167.0	Air Canada	YXX	
122	Aberdeen	United Kingdom	186.0	KLM	ABZ	
526	Abu Dhabi	United Arab Emirates	510.0	KLM	AUH	
8	Abuja	Nigeria	790.0	Turkish Airlines	ABV	
307	Acapulco	Mexico	3603.0	Multiple Airlines	ACA	

	Date	Link
295	2020-07-15	<a href="http://kayak.com/flights/TXL-YXX/2020-07-15/20...">http://kayak.com/flights/TXL-YXX/2020-07-15/20...</a>
122	2020-07-02	<a href="http://kayak.com/flights/TXL-ABZ/2020-07-02/20...">http://kayak.com/flights/TXL-ABZ/2020-07-02/20...</a>
526	2020-07-02	<a href="http://kayak.com/flights/TXL-AUH/2020-07-02/20...">http://kayak.com/flights/TXL-AUH/2020-07-02/20...</a>
8	2020-07-02	<a href="http://kayak.com/flights/TXL-ABV/2020-07-02/20...">http://kayak.com/flights/TXL-ABV/2020-07-02/20...</a>
307	2020-07-02	<a href="http://kayak.com/flights/TXL-ACA/2020-07-02/20...">http://kayak.com/flights/TXL-ACA/2020-07-02/20...</a>

We have about 500 rows of flight information, so we really need a good way to summarize the data. Let's return a dataframe showing the cheapest flights from our destination to each continent. Unfortunately, the JSON string returned by Kayak doesn't specify the continent of each destination country, so let's scrape wikipedia to map countries to their respective continents.

```

In [10]: def scrape_wikipedia():
        """

```

*This function scrapes and parses several wikipedia pages to map flight destination cities to their respective continents. It then cleans the resulting data frame.*

*Returns:*

*df - a data frame containing all cities from the wiki data and the continents where they are located*  
 """

```

urls = ['https://en.wikipedia.org/wiki/List_of_African_countries_by_area',
        'https://en.wikipedia.org/wiki/List_of_North_American_countries_by_GDP_(nominal)',
        'https://en.wikipedia.org/wiki/List_of_South_American_countries_by_population',
        'https://en.wikipedia.org/wiki/List_of_European_countries_by_area',
        'https://en.wikipedia.org/wiki/List_of_Oceanian_countries_by_population',
        'https://en.wikipedia.org/wiki/List_of_countries_in_Asia-Pacific_by_GDP_(nominal)',
        'https://en.wikipedia.org/wiki/List_of_Middle_Eastern_countries_by_population']

# We'll count the Middle East as Asia
continents = ['Africa', 'North America', 'South America', 'Europe',
              'Oceania', 'Asia', 'Asia']
all_continents, countries = [], []
df = pd.DataFrame(columns=['Country', 'Continent'])

for i in range(len(urls)):
    html = ur.urlopen(urls[i]).read()
    soup = BeautifulSoup(html, 'html.parser')
    table = soup.find_all('table')[0]
    rows = table.find_all('tr')

    for row in rows:
        columns = row.find_all('td')
        if len(columns) > 0:
            country = columns[1].get_text().strip()
            if country not in countries:
                countries.append(country)
                all_continents.append(continents[i])

# Remove all parentheses
countries = pd.Series(countries).replace(regex=True,
                                          to_replace=[r'\d', r'\([^)]*\)', ''], value=r'')

# Remove brackets and asterisks
countries = countries.replace(regex=True,
                              to_replace=[r'\[[^()]*\]', r'\[*]'],
                              value=r'')

df['Country'] = countries
df['Continent'] = pd.Series(all_continents)

return df

```

```

# If we've already run the scraper, there's no need to scrape wikipedia a
# second time.
if not os.path.isfile('data\continents.csv'):
    all_continents = scrape_wikipedia()
else:
    all_continents = pd.read_csv('data\continents.csv', index_col=0)

print(all_continents.head())

```

	Country	Continent
0	Algeria	Africa
1	Democratic Republic of the Congo	Africa
2	Sudan	Africa
3	Libya	Africa
4	Chad	Africa

Now that we have the continent for each country, let's find the best deal for each continent, and include two other regional destinations we are interested in to group alongside the continents (Japan and Hawaii).

```

In [3]: def summarize_results(cities):
        """
        This function finds the lowest priced flight to each continent, as well as
        to specific regions we're interested in, in this case Japan and Hawaii.

        Parameters:
        cities - a data frame with scraped kayak flight information with a
        continent mapped to each city.

        Returns:
        deals - a data frame containing flight information for the cheapest flight
        to each destination of interest.
        """

        hi_airports = ['HNL', 'MKK', 'OGG', 'KOA', 'ITO']
        hawaii = cities[cities['Airport'].str.match('LIH')]

        # Create a dataframe with all of the Hawaii flights
        for airport in hi_airports:
            hawaii = hawaii.append(cities[cities['Airport'].str.match(airport)])

        # Doing the same for Japan is a bit easier since we can just grep the
        # country ccolumn
        japan = cities[cities['Country'].str.match('Japan')]
        jp_lowest = japan.loc[japan['Price'].idxmin()].copy()
        jp_lowest['Continent'] = 'Japan*' # Differentiate the Japan flights from Asia flights

```

```

hi_lowest = hawaii.loc[hawaii['Price'].idxmin()].copy()
hi_lowest['Continent'] = 'Hawaii*'
lowest = cities.groupby(['Continent'])['Price'].idxmin()
deals = cities.iloc[lowest,:]
deals = deals.append(jp_lowest)
deals = deals.append(hi_lowest)
deals = deals.set_index('Continent')

return deals

flights_list = all_flights.merge(all_continents, how='left', on='Country')
results = summarize_results(flights_list)
print(results.head())

```

	City	Country	Price	Airline	Airport	\
Continent						
Africa	Marrakesh	Morocco	162.0	TAP AIR	PORTUGAL	RAK
Asia	Izmir	Turkey	150.0	Ukraine Intl	Air	ADB
Europe	Venice	Italy	69.0	Ryanair		TSF
North America	New York	United States	309.0	Multiple Airlines		JFK
Oceania	Perth	Australia	571.0	Scoot		PER

	Date	Link
Continent		
Africa	2020-07-17	<a href="http://kayak.com/flights/TXL-RAK/2020-07-17/20...">http://kayak.com/flights/TXL-RAK/2020-07-17/20...</a>
Asia	2020-07-19	<a href="http://kayak.com/flights/TXL-ADB/2020-07-19/20...">http://kayak.com/flights/TXL-ADB/2020-07-19/20...</a>
Europe	2020-07-15	<a href="http://kayak.com/flights/SXF-TSF/2020-07-15/20...">http://kayak.com/flights/SXF-TSF/2020-07-15/20...</a>
North America	2020-06-02	<a href="http://kayak.com/flights/TXL-JFK/2020-06-02/20...">http://kayak.com/flights/TXL-JFK/2020-06-02/20...</a>
Oceania	2020-08-12	<a href="http://kayak.com/flights/TXL-PER/2020-08-12/20...">http://kayak.com/flights/TXL-PER/2020-08-12/20...</a>

Here is a function that determines whether or not to send an email based on the present results.

```

In [13]: def check_df(results, start, end):
        """
        This function compares the results of the current scrape with previous
        results to determine if an email update should be sent.

        Parameters:
        results, start, end - dataframe with summarized scrape results, integer
        representing earliest possible departure date in YYYYMMDD format, integer
        representing latest return date

        Returns:
        msgs, email - list of strings indicating continents for which good deals are
        available, boolean indicating whether an email should be sent
        """

```

```

filename = 'data/' + str(start) + '_to_' + str(end) + '_kayak_scrape.csv'
if os.path.isfile(filename):
    df = pd.read_csv(filename)
else:
    df = pd.DataFrame(columns=['Date'])

current_scrape = results['Price']

# Append the current scrape as a row if it isn't a duplicate
if df.append(current_scrape).drop([
    'Date'], axis=1).duplicated().any() == False:
    df = df.append(current_scrape)
    now = datetime.datetime.now()
    df.iloc[-1,0] = now
    df.index = range(len(df))

df.to_csv(filename, index=False)
msgs = ''
# Bool indicating if an email will be sent. Will be set to true if good
# deals are detected
email = False
# Percent of average flight price to a given destination; if a current price is
# under this threshold, an email alert will be generated
email_threshold = 0.85

for column in range(1,len(df.columns)):
    col_mean = df.iloc[:,column].mean()
    if col_mean * email_threshold > df.iloc[-1,column]:
        name = df.columns[column]
        msg = 'Flights to ' + name + ' right now are abnormally cheap.<br>'
        msgs += msg
        email = True

print(df)

return msgs, email

msgs, email = check_df(results, timeframe_begin, timeframe_end)
print(msgs)

```

	Date	Africa	Asia	Europe	Hawaii*	Japan*	\
0	2019-09-06 21:18:53.349374	149	330	63	1096	525	
1	2019-09-06 21:26:08.604121	149	330	63	1096	525	
2	2019-09-06 21:40:52.898861	149	330	63	1096	525	
3	2019-09-08 09:52:53.051645	195	128	58	1160	728	
4	2019-09-09 12:33:22.396911	170	146	64	1441	844	
5	2019-09-09 17:31:59.685267	162	144	62	1197	547	
6	2019-09-10 09:36:00.039766	159	128	62	1306	618	

7	2019-09-10 17:09:53.544494	253	162	65	1197	618
8	2019-09-10 21:09:37.573083	253	133	65	1197	496
9	2019-09-11 10:29:23.825930	253	154	65	1059	631
10	2019-09-11 21:31:41.159486	170	128	65	1059	529
11	2019-09-11 22:20:12.757514	170	128	65	1059	529
12	2019-09-12 09:09:22.146991	170	128	65	1059	529
13	2019-09-12 11:47:40.541389	170	128	65	1059	558
14	2019-09-12 14:45:01.674894	170	329	65	1059	558
15	2019-09-13 18:44:08.345119	161	133	61	1059	517
16	2019-09-14 07:50:12.079461	161	128	61	1324	517
17	2019-09-14 08:57:42.802206	161	128	61	1324	517
18	2019-09-14 14:58:59.527522	193	133	61	1196	642
19	2019-09-15 10:26:40.968860	165	145	63	1326	541
20	2019-09-15 22:31:52.069036	165	125	61	1054	541
21	2019-09-16 09:49:09.023443	165	125	66	1326	566
22	2019-09-16 12:14:03.513101	165	125	66	1326	566
23	2019-09-16 20:25:28.076240	162	137	66	1196	546
24	2019-09-17 09:49:48.161516	162	138	70	1326	545
25	2019-09-17 13:57:56.650502	162	150	69	1059	545

	North America	Oceania	Oklahoma	South America
0	366	698	563.0	579
1	366	698	520.0	579
2	366	698	551.0	579
3	477	745	544.0	584
4	364	697	551.0	624
5	364	618	552.0	624
6	485	733	531.0	624
7	368	691	560.0	598
8	368	691	560.0	598
9	396	691	567.0	628
10	311	698	282.0	593
11	304	698	282.0	593
12	312	698	557.0	666
13	311	697	539.0	666
14	311	697	529.0	666
15	282	697	523.0	584
16	309	805	525.0	628
17	309	805	519.0	628
18	296	700	525.0	598
19	315	800	519.0	598
20	309	703	525.0	598
21	301	681	519.0	619
22	301	681	NaN	598
23	311	572	NaN	598
24	309	572	529.0	598
25	309	571	NaN	598

Flights to Oceania right now are abnormally cheap.<br>

This function sends an email if the current scrape has any outstanding deals to at least one of our destinations.

```
In [14]: def send_email(flights):
        """
        This function sends an email with the summarized flight data as a data
        frame in html to the specified address.

        Parameters:
        flights - a dataframe of the best deals on flights and the corresponding
        details returned by our scrape.
        """

        password = input('Type your password:')
        message = MIMEMultipart('alternative')
        message.add_header('Content-Type', 'html')
        sender = 'youremail@domain.com'
        receiver = 'theiremail@domain.com'
        message['Subject'] = "Here is your latest Kayak scrape!"
        message['From'] = sender
        message['To'] = receiver
        html = "<html><head></head><body><p>" + msgs + \
            flights.to_html() + ". </p></body></html>"

        part1 = MIMEText(html, 'html')
        message.attach(part1)
        mail = smtplib.SMTP('smtp.gmail.com', 587)
        mail.ehlo()
        mail.starttls()
        mail.login('your_username', password)
        mail.sendmail(sender, receiver.split(','), message.as_string())
        mail.quit()

    if email:
        #send_email(results)
        pass # No email in this jupyter version
```

This function saves a dataframe with the data from all flights (not just the best deals), which is useful for producing some charts.

```
In [17]: def save_scrape(flights, start, end):
        """
        This function appends a column containing all of the current scraped prices
        to a dataframe of saved flight price data.

        Parameters:
        flights, start, end - dataframe with all scraped flight results, integer
```



*representing earliest possible departure date in YYYYMMDD format, integer  
representing latest return date*

*Returns:*

*dataframe containing all previously scraped flight data plus the current scrape  
"""*

```
filename = 'data/' + str(start) + '_to_' + str(end) + '_all_flights.csv'
now = datetime.datetime.now()
current_prices = flights.set_index('City')['Price']
current_prices.name = now
```

```
if os.path.isfile(filename):
    df = pd.read_csv(filename, index_col=0)
    df = df.merge(current_prices, how='outer', left_index=True,
                  right_index=True)
else:
    df = pd.DataFrame(current_prices)
    df['Continent'] = flights_list.set_index('City')['Continent']

df.to_csv(filename)

return df.drop('Continent', axis=1).dropna()
```

```
historical = save_scrape(flights_list, timeframe_begin, timeframe_end)
print(historical.head())
```

```
2019-09-11 22:05:24.869258  2019-09-12 09:01:14.631851  \
City
Aberdeen                563.0                188.0
Abu Dhabi                510.0                510.0
Abuja                  1877.0                1878.0
Acapulco               2486.0                2486.0
Accra                   711.0                711.0
```

```
2019-09-12 11:36:55.827211  2019-09-12 12:22:00.953972  \
City
Aberdeen                188.0                188.0
Abu Dhabi                510.0                510.0
Abuja                  1878.0                1878.0
Acapulco               2486.0                2486.0
Accra                   709.0                709.0
```

```
2019-09-12 12:23:51.262758  2019-09-12 12:47:19.909551  \
City
Aberdeen                188.0                188.0
Abu Dhabi                510.0                510.0
Abuja                  1878.0                1878.0
Acapulco               2486.0                2486.0
```

Accra	709.0	709.0	
	2019-09-12 14:41:09.103147	2019-09-12 14:43:29.661475	\
City			
Aberdeen	679.0	679.0	
Abu Dhabi	510.0	510.0	
Abuja	1878.0	1878.0	
Acapulco	3591.0	3591.0	
Accra	669.0	669.0	
	2019-09-13 18:31:06.671201	2019-09-13 18:33:38.926318	... \
City			...
Aberdeen	185.0	185.0	...
Abu Dhabi	499.0	499.0	...
Abuja	1874.0	1874.0	...
Acapulco	3584.0	3584.0	...
Accra	710.0	710.0	...
	2019-09-15 22:29:49.689412	2019-09-16 09:47:06.249882	\
City			
Aberdeen	186.0	186.0	
Abu Dhabi	502.0	502.0	
Abuja	1876.0	790.0	
Acapulco	3587.0	3587.0	
Accra	643.0	643.0	
	2019-09-16 12:14:03.140103	2019-09-16 20:11:48.235263	\
City			
Aberdeen	186.0	186.0	
Abu Dhabi	502.0	510.0	
Abuja	790.0	790.0	
Acapulco	3587.0	3587.0	
Accra	465.0	465.0	
	2019-09-16 20:27:21.839269	2019-09-17 09:44:16.044334	\
City			
Aberdeen	186.0	186.0	
Abu Dhabi	510.0	510.0	
Abuja	790.0	790.0	
Acapulco	3587.0	3587.0	
Accra	465.0	465.0	
	2019-09-17 09:47:57.780076	2019-09-17 14:01:20.505055	\
City			
Aberdeen	186.0	186.0	
Abu Dhabi	510.0	510.0	
Abuja	790.0	790.0	
Acapulco	3587.0	3603.0	

Accra	465.0	717.0
	2019-09-17 14:01:25.940395	2019-09-17 14:01:55.559245
City		
Aberdeen	186.0	186.0
Abu Dhabi	510.0	510.0
Abuja	790.0	790.0
Acapulco	3603.0	3603.0
Accra	717.0	717.0

[5 rows x 29 columns]

This chart summarizes the distribution of prices to each continent from the current scrape.

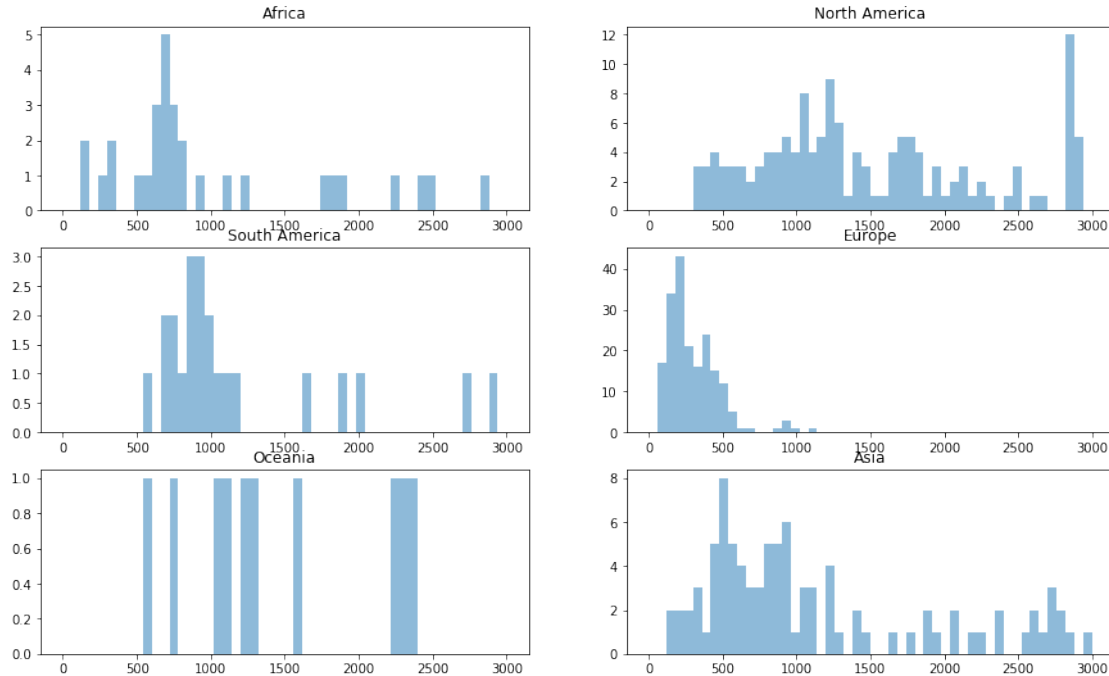
```
In [22]: %matplotlib inline
          %pylab inline
          pylab.rcParams['figure.figsize'] = (15, 9)

          continents = ['Africa', 'North America', 'South America',
                        'Europe', 'Oceania', 'Asia']
          plot_num = 1

          for continent in continents:
              plt.subplot(3,2,plot_num)
              df = flights_list.where(flights_list['Continent']==continent).dropna()
              plt.hist(df['Price'], bins=50, alpha=0.5, range=(0,3000))
              title = plt.gca().set_title(continent)
              plot_num += 1

          plt.show()
```

Populating the interactive namespace from numpy and matplotlib



The next chart is possibly more useful, because it tells us which cities in our regional destinations have a wide variability in flight prices. Cities for which the boxes are long or have many outliers may be worth paying additional attention to.

```
In [21]: dest_prices, labels = [], []
```

```
# This lists cities from Hawaii and Japan to which flights are present in our
# data frame.
bp_cities = ['Honolulu', 'Kailua-Kona', 'Kahului', 'Lihue', 'Osaka', 'Nagoya',
             'Tokyo', 'Sapporo', 'Okinawa']

for city in bp_cities:
    df = historical[historical.index.str.match(city)].T
    if df.shape[1] > 0:
        df = pd.Series(df[city]).tolist()
        dest_prices.append(df)
        labels.append(city)

fig, ax = plt.subplots()
ax.boxplot(dest_prices, labels=labels, whis=2)
plt.ylabel('Price in USD')
plt.show()
```

